

ANLERNEN VON NETZWERK-TRAFFIC

Learning & Soft Computing

Learning & Soft Computing
ANLERNEN VON NETZWERK-TRAFFIC

Nicola Jäger, Marc Schmidt, Jonathan Ebinger

Agenda





- Datenzugriff über Google Drive
- Zusammenarbeiten begrenzt möglich
- Pro Ressourcen für High-Performance RAM
- GPU Beschleunigung



- Probleme mit Internetzugriff
- Sehr langsame Ausführungszeiten
- Datenzugriff nicht weiter ausprobiert



Jupyter Notebook

- **Docker Hub:**
jupyter/scipy-notebook
- **Daten befinden sich in einem persistenten Datenträger**
- **Zusammenarbeiten besser als bei colab, aber dennoch nicht optimal**
- **Ausführungszeiten fast identisch mit Google colab**
- **24/7 Uptime**
- **Keine Kosten**



DataSpell

- Dateien lokal auf OS
- Schnellere Ausführungszeiten
- Kein Zusammenarbeiten möglich



Python

- Dateien lokal auf OS
- Sehr schnelle Ausführungszeiten
- Kein Zusammenarbeiten möglich

Datensatz: **Aposemat IoT-23**

- Erstellung im Rahmen des Avast AIC-Labors mit finanzieller Unterstützung von Avast
- Im Januar 2020 veröffentlicht, umfasst Aufnahmen aus den Jahren 2018 bis 2019
- Beinhaltet echten Netzwerkverkehr von IoT-Geräten
 - Infizierter / böstiger IoT-Netzwerk-Traffic
 - Mirai Botnet
 - Okiru Botnet
 - PortScan
 - C&C-HeartBeat
 - DDoS
 - unbedenklicher/normaler IoT-Netzwerk-Traffic

Vorstellung der IDS-Datensätze

- Name der Datensätze: CTU-IoT-Malware-*

Datensatz	Normal Traffic Flows	Malicious Traffic Flows
Capture-1-1 (Hide and Seek) [1]	469.275	539.465
Capture-7-1 (Linux.Mirai) [2]	75.955	11.378.759
Capture-35-1 (Mirai) [3]	8.262.389	2.185.398

Malicious Flows bestehen teilweise aus weiteren Kategorien, wie:

- C&C
- C&C-HeartBeat
- C&C-FileDownload
- DDoS
- PartOfAHorizontalPortScan

Raw .log files

ts	uid	orig_h	orig_p	resp_h	resp_p	proto	service	duration	orig_bytes
1.525880e+09	CUmrqr4svHuSXJy5z7	192.168.100.103	51524	65.127.233.163	23	tcp	NaN	2.999051	0

resp_bytes	conn_state	local_orig	local_resp	missed_bytes	history	orig_pkts
0	S0	NaN	NaN	0	S	3

orig_ip_bytes	resp_pkts	resp_ip_bytes	label	detailed-label
180	0	0	Malicious	PartOfAHorizontalPortScan

.log → .csv

orig_h	orig_p	resp_h	resp_p	proto
192.168.100.103	51524	65.127.233.163	23	Y

conn_state	missed_bytes	history	orig_pkts
S0	0	X	3

orig_ip_bytes	resp_pkts	resp_ip_bytes	label
180	0	0	Malicious

Encoding von:

- IP-Adresse
- History
- Protocol
- Label

Zusammenführen der Dateien

1. Datensatz 1 - Datei 1
 - Verteilung: 500.000 : 400.000

2. Datensatz merged - Datei 1 + 7 + 37
 - Nach Konvertierung, Verteilung: 1.000.000 : 1.000.000

Split der Daten in Train/Test/Valid

- Aufspalten in Train **60%**, Test **20%** und Valid **20%** Datensätze
- Transformation der Daten durch **Standardisierung**
- Berechnung von **Mittelwert und Standardabweichung**, um Merkmale oder Variablen in einem Datensatz auf vergleichbare Skalen zu bringen
- Zufälliges **Oversampling** der Train Datensätze (Gleichverteilen)

```
1 def split(df):
2     train, valid, test = np.split(df.sample(frac=1), [int(0.6*len(df)), int(0.8*len(df))])
3     return train, valid, test
4
5 def scale(df: pd.DataFrame, oversample = False):
6     X = df.drop(['label'], axis=1)
7     y = df['label'].values
8
9     scaler = StandardScaler()
10    X = scaler.fit_transform(X)
11    if oversample:
12        ros = RandomOverSampler()
13        X, y = ros.fit_resample(X, y)
14
15    data = np.hstack((X, np.reshape(y, (-1, 1))))
16    return data, X, y
17
18 train, valid, test = split(SOURCEDF)
19 train, X_train, y_train = scale(train, True)
20 valid, X_valid, y_valid = scale(valid, False)
21 test, X_test, y_test = scale(test, False)
```

Sneak-Peek in die Daten

- Daten wurden **standardisiert**
- 60% - _train
- 20% - _test
- 20% - _vaild

Die Datensätze werden zufällig erstellt
-> Genauigkeiten, können sich nach
Datensatz unterscheiden



```
1 pd.DataFrame(X_train)
2 """
3
4      0      1      2      3      4
5 0 0.092973 -0.070335 1.792180 1.711901 1.178841
6 1 0.092973 -0.470168 -1.720678 -0.819821 -0.740578
7 2 0.092973 -1.093452 0.839101 -0.702064 -0.740578
8 3 0.092973 1.147476 1.240883 -0.333229 -0.740578
9 4 0.092973 -0.070335 -1.031813 1.306920 1.178841
10
11 """
```



```
1 pd.Series(y_train).head()
2 """
3 0 0
4 1 1
5 2 1
6 3 1
7 4 0
8 Name: label, dtype: int64
9 """
```

Datensatz 1

kNN - K-Optimierung

Die idealen Parameter wurden im Vorfeld durch ein Ausprobieren ermittelt, später durch GridSearchCV

```
1 k_range=range(1,10)
2 knn_r_acc = []
3
4 for k in k_range:
5     knn = KNeighborsClassifier(n_neighbors=k, algorithm='ball_tree', weights='distance', metric='euclidean', n_jobs=-1)
6     knn.fit(X_train,y_train)
7
8     train_score = knn.score(X_train,y_train)
9     test_score = knn.score(X_test,y_test)
10    vaild_score = knn.score(X_valid,y_valid)
11
12    knn_r_acc.append((k, train_score,test_score, vaild_score))
13
14    print(f"finished {k}")
15
16 df = pd.DataFrame(knn_r_acc, columns=['K', 'Train Score', 'Test Score', 'Vaild Score'])
17 print(df)
```

Laufzeit: 15 min

k = 2

- Genauigkeit von 99%
- **Achtung: Werte sind gerundet**



1	K	Train Score	Test Score	Vaild Score
2	1	1.0	0.998721	0.998533
3	2	1.0	0.998721	0.998533 #
4	3	1.0	0.998374	0.998171
5	4	1.0	0.998141	0.997898
6	5	1.0	0.997834	0.997398
7	6	1.0	0.997611	0.997120
8	7	1.0	0.997279	0.996798
9	8	1.0	0.996976	0.996506
10	9	1.0	0.996778	0.996312



```
1 knn = KNeighborsClassifier(n_neighbors=2, algorithm='ball_tree',
2   weights='distance', metric='euclidean', n_jobs=-1)
3 knn.fit(X_train,y_train)
4 y_pred = knn.predict(X_test)
5
6 print(classification_report(y_test, y_pred))
7
8 ...
9
10
11
12
13
14
15
16
17
```

		precision	recall	f1-score	support
	0	1.00	1.00	1.00	94155
	1	1.00	1.00	1.00	107594
	accuracy			1.00	201749
	macro avg	1.00	1.00	1.00	201749
	weighted avg	1.00	1.00	1.00	201749
	...				

Laufzeit: 2 min

Nicola

<https://10015.io/tools/code-to-image-converter>

python, theme: AtomOne Light, Settings Line Number on, Rest default

no background

compact

dTrees - Hyperparameter Tuning



```
1 def dtree_grid_search(X,y,nfolds):
2     #create a dictionary of all values we want to test
3     param_grid = { 'ccp_alpha': np.arange(0,0.8,0.01), 'max_depth': [1,2,3,4], 'random_state': [1,2,3,4,5], 'criterion': ['entropy', 'gini']}
4     # decision tree model
5     dtree_model=DecisionTreeClassifier()
6     #use gridsearch to test all values
7     dtree_gscv = GridSearchCV(dtree_model, param_grid, cv=nfolds, n_jobs=30)
8     #fit model to data
9     dtree_gscv.fit(X, y)
10    # print best score achieved
11    print(dtree_gscv.best_score_)
12
13    return dtree_gscv.best_params_
```

Laufzeit: ca. 5 min (auf 48 Kernen)



```
1 0.9935989770664833
2 {'ccp_alpha': 0.0, 'criterion': 'gini', 'max_depth': 4, 'random_state': 1}
```

dTrees - Hyperparameter Tuning

```
1  # Create Decision Tree classifier object
2  dtree = DecisionTreeClassifier(max_depth=4, random_state=1, criterion= "gini")
3
4  # Train Decision Tree Classifier
5  dtree = dtree.fit(X_train,y_train)
6
7  dtree.get_params()
```

Laufzeit: 3 sek

```
1  {'ccp_alpha': 0,
2   'class_weight': None,
3   'criterion': 'gini',
4   'max_depth': 4,
5   'max_features': None,
6   'max_leaf_nodes': None,
7   'min_impurity_decrease': 0.0,
8   'min_samples_leaf': 1,
9   'min_samples_split': 2,
10  'min_weight_fraction_leaf': 0.0,
11  'random_state': 1,
12  'splitter': 'best'}
```


dTrees - Classification report



```
1 # Predict the response for validation dataset
2 y_pred = dtree.predict(X_valid)
3
4 # Create classification report
5 classification_report(y_valid, y_pred, target_names=['benign', 'malicious'])
```




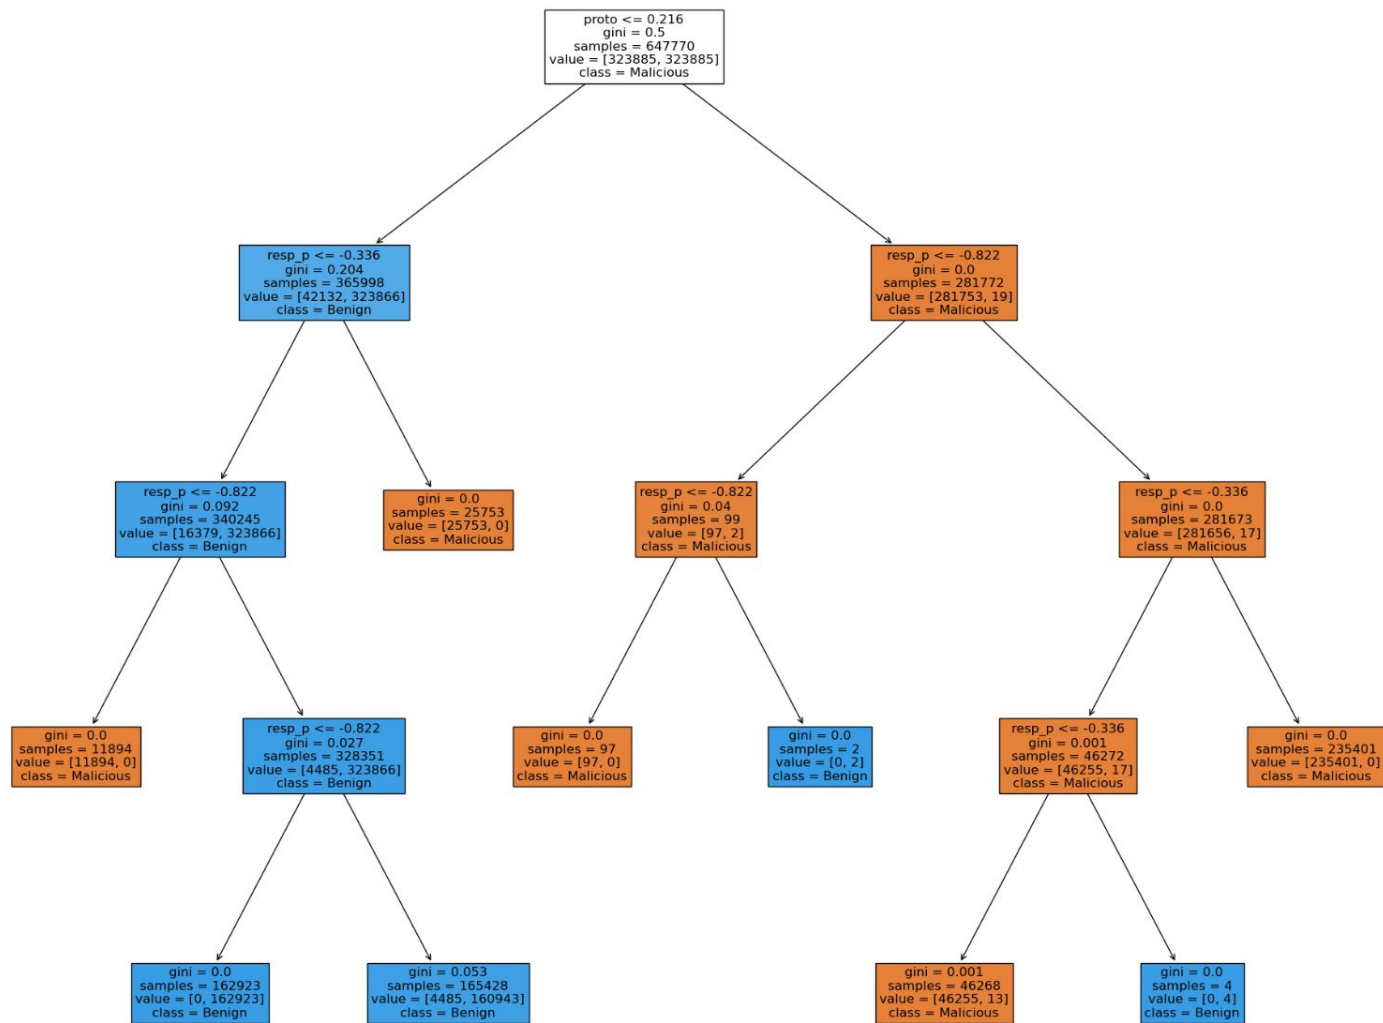
1		precision	recall	f1-score	support
2					
3	benign	1.00	0.99	0.99	94141
4	malicious	0.99	1.00	0.99	107608
5					
6	accuracy			0.99	201749
7	macro avg	0.99	0.99	0.99	201749
8	weighted avg	0.99	0.99	0.99	201749

dTrees - Feature importances




```
1 df_ds1_drop = df_ds1.drop(['label'], axis=1)
2 feature_names = df_ds1_drop.columns
3 feature_importance = pd.DataFrame(dtrees.feature_importances_, index = feature_names).sort_values(0, ascending=False)
4
5 feature_importance.head(10).plot(kind='bar')
```





aNN

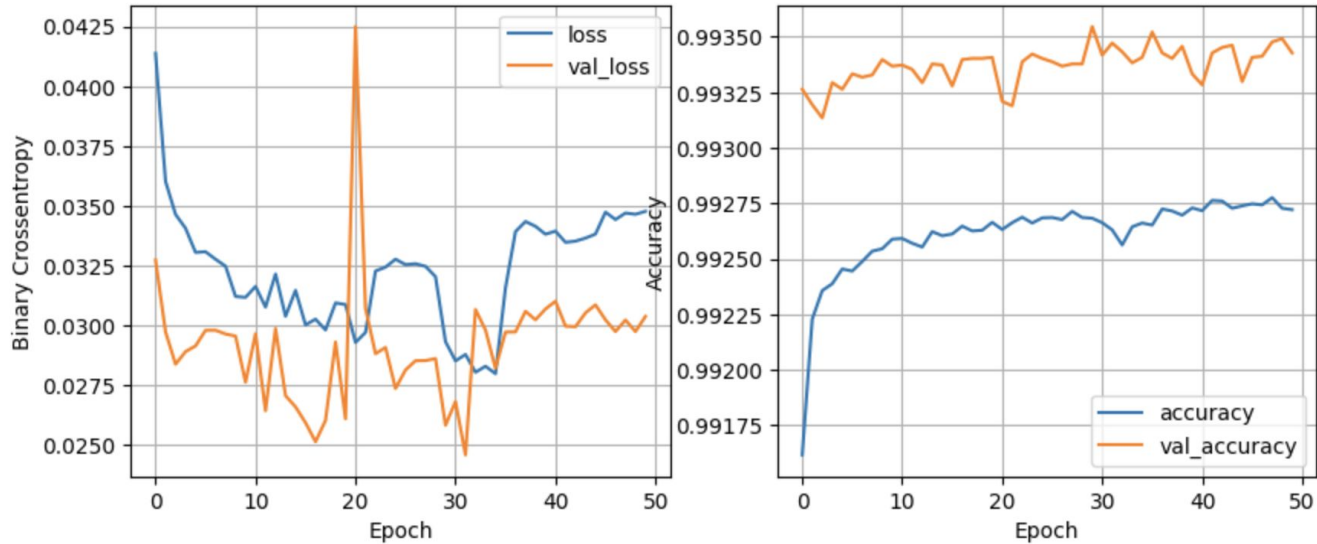




```
1 nn_model = tf.keras.Sequential([
2     tf.keras.layers.Dense(nodes, activation='relu', input_shape=(5,)),
3     # take certain nodes at a specific rate, and dont train them
4     tf.keras.layers.Dropout(dropout_prob),
5     tf.keras.layers.Dense(nodes, activation='relu'),
6     tf.keras.layers.Dropout(dropout_prob),
7     # this will set the output to 0 or 1, which helps with the classification
8     tf.keras.layers.Dense(1, activation='sigmoid')
9 ])
```



```
1 for nodes in [16, 32, 64]:
2     for dropout_prob in [0, 0.2]:
3         for lr in [0.1, 0.005, 0.001]:
4             for batch_size in [32, 64, 128]:
5                 print(f"{nodes} nodes, dropout_prob {dropout_prob}, lr {lr}, batch size {batch_size}")
6                 model, history = nn_train(X_train=X_train, y_train=y_train, X_valid=X_valid, y_valid=y_valid,
7                                           epochs=epochs, nodes=nodes, dropout_prob=dropout_prob, lr=lr, batch_size=batch_size)
8                 plot_nn(history=history)
9                 val_loss = model.evaluate(X_test, y_test)[0]
10                if val_loss < least_val_loss:
11                    least_val_loss = val_loss
12                    least_loss_model = model
```



6305/6305 [=====] - 5s 760us/step - loss: 0.0304 - accuracy: 0.9934
32 nodes, dropout_prob 0.2, lr 0.001, batch size 32

Laufzeit: 25 min

aNN - Gridsearch

```
1 def create_model(neurons=1, activation='relu', dropout_prob=0):
2     # create model
3     ann = Sequential()
4     ann.add(Dense(units = neurons, activation = activation, input_shape=(5,)))
5     ann.add(Dense(units = 1))
6     ann.add(Dropout(dropout_prob))
7     # Compile model
8     ann.compile(optimizer = 'adam', loss = 'mean_squared_error')
9     return ann
10 # fix random seed for reproducibility
11 np.random.seed(0)
12 # create model
13 ann = KerasRegressor(build_fn = create_model, epochs = 100, batch_size = 10, verbose = 0)
14 # define the grid search parameters
15 parameters = {'neurons': [5, 10, 15, 20, 25, 30, 35, 40, 45, 50, 55, 60],
16               'activation': ['softmax', 'softplus', 'softsign', 'relu', 'tanh', 'sigmoid', 'hard_sigmoid', 'linear'],
17               'dropout_prob': [0,0.2]}
18 grid = GridSearchCV(estimator = ann, param_grid = parameters, n_jobs = -1, cv=3)
19 grid_result = grid.fit(X_train, y_train)
20 # summarize results
21 print(f"Best: {grid_result.best_score_} using {grid_result.best_params_}")
22 means = grid_result.cv_results_['mean_test_score']
23 stds = grid_result.cv_results_['std_test_score']
24 params = grid_result.cv_results_['params']
```

Abbruch nach 14 st

Datensatz merged

kNN - K-Optimierung (Zusammengefügter Datensatz)



```
1 X = SOURCEDF.drop('label',axis=1)
2 y = SOURCEDF.label
3
4 from sklearn.model_selection import GridSearchCV
5
6 grid_params = {
7     'n_neighbors': [1,2,3,4,5,6,7],
8     'weights': ['uniform', 'distance'],
9     'metric': ['euclidean', 'manhattan']
10 }
11
12 gs = GridSearchCV(
13     KNeighborsClassifier(),
14     grid_params,
15     verbose = 1,
16     cv = 3,      # Determines the cross-validation splitting strategy
17                 # integer, to specify the number of folds
18     n_jobs = -1 # use all processors to perform the model
19 )
20
21 gs_results = gs.fit(X, y)
```

Ursprungsdaten

Fitting 3 folds for each of 28
candidates, totalling 84 fits

Laufzeit: 5 min

kNN - K-Optimierung (Zusammengefügter Datensatz)



```
1 gs_results.best_score_  
2 0.8986799842095562  
3  
4 gs_results.best_params_  
5 {'metric': 'manhattan', 'n_neighbors': 2, 'weights': 'distance'}
```


kNN - K-Optimierung (Zusammengefügter Datensatz)

Laufzeit: 2 min

```
1 knn = KNeighborsClassifier(n_neighbors=2, metric='manhattan', weights='distance', n_jobs=-1)
2 knn.fit(X_train,y_train)
3
4 knn.score(X_train,y_train)
5 1.0
6
7 knn.score(X_test,y_test)
8 0.897141417305034
9
10 knn.score(X_valid,y_valid)
11 0.9162863534675615
```

```
1 y_pred = knn.predict(X_valid)
2 print(classification_report(y_valid, y_pred))
3
4 ...
5
6
7 precision recall f1-score support
8
9
10 0 0.93 0.90 0.92 212630
11 1 0.91 0.93 0.92 212020
12
13
14 accuracy 0.92 424650
15 macro avg 0.92 0.92 0.92 424650
16 weighted avg 0.92 0.92 0.92 424650
17 ...
```

kNN - K-Optimierung (Zusammengefügter Datensatz)



kNN - GridSearchCV mit mehr Parametern

'weights': ['uniform', 'distance'],
'metric': ['euclidean', 'manhattan', '**'haversine'**,
'nan_euclidean'],

Was macht man, wenn man nicht genug
CPU/RAM Leistung hat?

-> man nutzt die AWS Cloud

Model	vCPU	Memory (GiB)
c5.large	2	4
c5.xlarge	4	8
c5.2xlarge	8	16
c5.4xlarge	16	32
c5.9xlarge	36	72

```
0[ 100.0%] 4[ 100.0%] 8[ 100.0%] 12[ 100.0%]
1[ 100.0%] 5[ 100.0%] 9[ 100.0%] 13[ 100.0%]
2[ 100.0%] 6[ 100.0%] 10[ 100.0%] 14[ 100.0%]
3[ 100.0%] 7[ 100.0%] 11[ 100.0%] 15[ 100.0%]
Mem[ 24.3G/30.6G] Tasks: 60, 58 thr, 171 kthr; 16 running
Swp[ 0K/0K] Load average: 4.54 4.34 2.20
Uptime: 00:11:37

Main I/O
PID USER PRI NI VIRT RES SHR S CPU%MEM% TIME+ Command
26786 ec2-user 20 0 1851M 1494M 158M R 100.6 4.8 0:14.75 /usr/bin/python3 -m joblib.externals.loky.backend.popen loky_posix --process-name LokyProcess-4 --pipe 16
26787 ec2-user 20 0 1851M 1494M 158M R 100.6 4.8 0:14.73 /usr/bin/python3 -m joblib.externals.loky.backend.popen loky_posix --process-name LokyProcess-5 --pipe 17
26794 ec2-user 20 0 1867M 1509M 158M R 100.6 4.8 0:14.74 /usr/bin/python3 -m joblib.externals.loky.backend.popen loky_posix --process-name LokyProcess-12 --pipe 24
26795 ec2-user 20 0 1856M 1499M 159M R 100.6 4.8 0:14.73 /usr/bin/python3 -m joblib.externals.loky.backend.popen loky_posix --process-name LokyProcess-13 --pipe 25
26797 ec2-user 20 0 2870M 1553M 158M R 100.6 5.0 0:14.75 /usr/bin/python3 -m joblib.externals.loky.backend.popen loky_posix --process-name LokyProcess-15 --pipe 27
26798 ec2-user 20 0 1960M 1603M 159M R 100.6 5.1 0:14.75 /usr/bin/python3 -m joblib.externals.loky.backend.popen loky_posix --process-name LokyProcess-16 --pipe 28
26783 ec2-user 20 0 2867M 2262M 159M R 100.0 7.2 0:14.72 /usr/bin/python3 -m joblib.externals.loky.backend.popen loky_posix --process-name LokyProcess-1 --pipe 13
26788 ec2-user 20 0 1855M 1498M 158M R 100.0 4.8 0:14.74 /usr/bin/python3 -m joblib.externals.loky.backend.popen loky_posix --process-name LokyProcess-6 --pipe 18
26793 ec2-user 20 0 1852M 1495M 158M R 100.0 4.8 0:14.69 /usr/bin/python3 -m joblib.externals.loky.backend.popen loky_posix --process-name LokyProcess-11 --pipe 23
26796 ec2-user 20 0 1852M 1495M 159M R 100.0 4.8 0:14.75 /usr/bin/python3 -m joblib.externals.loky.backend.popen loky_posix --process-name LokyProcess-14 --pipe 26
26785 ec2-user 20 0 2867M 2500M 159M R 99.4 8.0 0:14.70 /usr/bin/python3 -m joblib.externals.loky.backend.popen loky_posix --process-name LokyProcess-3 --pipe 15
26790 ec2-user 20 0 2866M 1625M 159M R 99.4 5.2 0:14.76 /usr/bin/python3 -m joblib.externals.loky.backend.popen loky_posix --process-name LokyProcess-8 --pipe 20
26791 ec2-user 20 0 1851M 1494M 159M R 99.4 4.8 0:14.67 /usr/bin/python3 -m joblib.externals.loky.backend.popen loky_posix --process-name LokyProcess-9 --pipe 21
26792 ec2-user 20 0 1861M 1504M 158M R 99.4 4.8 0:14.75 /usr/bin/python3 -m joblib.externals.loky.backend.popen loky_posix --process-name LokyProcess-10 --pipe 22
26789 ec2-user 20 0 2866M 1894M 159M R 98.7 6.0 0:14.57 /usr/bin/python3 -m joblib.externals.loky.backend.popen loky_posix --process-name LokyProcess-7 --pipe 19
26784 ec2-user 20 0 2876M 1828M 159M R 97.4 5.8 0:14.68 /usr/bin/python3 -m joblib.externals.loky.backend.popen loky_posix --process-name LokyProcess-2 --pipe 14
```

kNN - K-Optimierung (Zusammengefügter Datensatz) GridSearchCV

Output

```
Fitting 3 folds for each of 40 candidates, totalling 120 fits
[CV 1/3; 1/40] START metric=euclidean, n_neighbors=1, weights=uniform.....
[CV 2/3; 1/40] START metric=euclidean, n_neighbors=1, weights=uniform.....
[CV 2/3; 2/40] START metric=euclidean, n_neighbors=1, weights=distance.....
[CV 3/3; 1/40] START metric=euclidean, n_neighbors=1, weights=uniform.....
[CV 2/3; 3/40] START metric=euclidean, n_neighbors=2, weights=uniform.....
[CV 1/3; 2/40] START metric=euclidean, n_neighbors=1, weights=distance.....
[CV 1/3; 3/40] START metric=euclidean, n_neighbors=2, weights=uniform.....
[CV 3/3; 3/40] START metric=euclidean, n_neighbors=2, weights=uniform.....
[CV 3/3; 2/40] START metric=euclidean, n_neighbors=1, weights=distance.....
[CV 1/3; 4/40] START metric=euclidean, n_neighbors=2, weights=distance.....
[CV 2/3; 4/40] START metric=euclidean, n_neighbors=2, weights=distance.....
[CV 2/3; 2/40] END metric=euclidean, n_neighbors=1, weights=distance;; score=0.908 total time= 9.8s
[CV 3/3; 4/40] START metric=euclidean, n_neighbors=2, weights=distance.....
[CV 2/3; 4/40] END metric=euclidean, n_neighbors=2, weights=distance;; score=0.908 total time= 16.0s
[CV 3/3; 2/40] END metric=euclidean, n_neighbors=1, weights=distance;; score=0.964 total time= 16.0s
[CV 1/3; 5/40] START metric=euclidean, n_neighbors=3, weights=uniform.....
```

Probleme:

- durch OOM-Killer
 - n_jobs = 11,
 - pre_dispatch = 11
- Sehr lange Ausführungszeiten! auch mit mehr Ressourcen
- Lösung: Maximal **zwei** Metrics verwendet, nicht alle gleichzeitig

Auszug aus den Ergebnissen:

```
'metric': ['l1', 'l2'] =  
'Algorithm': default
```

```
  {'metric': 'l1', 'n_neighbors': 2, 'weights': 'distance'}  
0.8986799842095562
```

```
'metric': ['euclidean', 'manhattan'],  
'algorithm': ['auto', 'ball_tree']
```

```
  {'algorithm': 'auto', 'metric': 'manhattan', 'n_neighbors': 2, 'weights': 'distance'}  
0.8986799842095562
```

Dennoch keine weiteren Verbesserungen auf den Split-Daten!

dTrees (Zusammengefügter Datensatz) - Hyperparameter Tuning



```
1 def dtree_grid_search(X,y,nfolds):
2     #create a dictionary of all values we want to test
3     param_grid = {'max_depth': [1,2,3,4], 'random_state': [None,1,2], 'criterion': ['entropy', 'gini']}
4     # decision tree model
5     dtree_model=DecisionTreeClassifier()
6     #use gridsearch to test all values
7     dtree_gscv = GridSearchCV(dtree_model, param_grid, cv=nfolds, n_jobs=30)
8     #fit model to data
9     dtree_gscv.fit(X, y)
10    # print best score achieved
11    print(dtree_gscv.best_score_)
12
13    return dtree_gscv.best_params_
```


Laufzeit: ca. 3 min (auf 48 Kernen)



```
1 0.9858344685226235
2 {'criterion': 'gini', 'max_depth': 4, 'random_state': 1}
```


dTrees (Zusammengefügter Datensatz)

Feature importances




Classification Report



1		precision	recall	f1-score	support
2					
3	benign	0.79	0.95	0.86	212148
4	malicious	0.94	0.74	0.83	212502
5					
6	accuracy			0.85	424650
7	macro avg	0.86	0.85	0.85	424650
8	weighted avg	0.86	0.85	0.85	424650


aNN (Zusammengefügter Datensatz)



loss: 0.2343 - accuracy: 0.8752

16 nodes, dropout_prob 0, lr 0.005, batch size 64

aNN (Zusammengefüger Datensatz)



loss: 0.3613 - accuracy: 0.7534
16 nodes, dropout_prob 0, lr 0.1, batch size 64

Vergleich der Trainings (beider Datensatz)

Modelle	Capture-1-1 Datensatz	Zusammengefügter Datensatz
kNN	99,85 %	91,62 %
dTrees	99,99 %	84,86 %
aNN	99,87 %	87,5%

INTRUSION DETECTOR LEARNING

<https://www.openml.org/search?type=data&status=active&sort=runs&id=42746>

HIKARI-2021: Generating Network Intrusion Detection Dataset Based on Real and Encrypted Synthetic Attack Traffic

<https://zenodo.org/record/5199540#.ZFh2FC9BxBo>

Kaggle IDS

<https://www.kaggle.com/datasets/amankumar255/network-intrusion-detection>

Cyber-security Datasets

<https://www.kaggle.com/discussions/general/335189>

Kyoto University's Honeypots

http://www.takakura.com/Kyoto_data/new_data201704/

IoT Dataset (das was wir nutzen)

<https://www.stratosphereips.org/datasets-iot23>

Dataset

Sebastian Garcia, Agustin Parmisano, & Maria Jose Erquiaga. (2020). IoT-23: A labeled dataset with malicious and benign IoT network traffic (1.0.0) [Data set]. Zenodo.

<https://doi.org/10.5281/zenodo.4743746>

[1]

<https://mcfp.felk.cvut.cz/publicDatasets/IoTDatasets/CTU-IoT-Malware-Capture-1-1/bro/conn.log.labeled>

[2]

<https://mcfp.felk.cvut.cz/publicDatasets/IoTDatasets/CTU-IoT-Malware-Capture-7-1/bro/conn.log.labeled>

[3]

<https://mcfp.felk.cvut.cz/publicDatasets/IoTDatasets/CTU-IoT-Malware-Capture-35-1/bro/conn.log.labeled>



Vielen Dank für Eure Aufmerksamkeit!